

Bucknell University

CHESS AI: A LOOK INTO AN ARTIFICIALLY INTELLIGENT CHESS AGENT

Final Project Report

Alex Godziela, Henry Kwan, Elias Strizower
CSCI 379 AI & Cognitive Science
Professor Chris Dancy

INTRODUCTION

Artificial Intelligence (AI) is a field of study that looks into intelligence exhibited by machines. Throughout our class, we have learned about the psychology behind AI and what it means for a robot to be intelligent. In order to explore AI more deeply, we decided to build a Chess AI for our final project. This intelligent computer system can optimize its moves in real-time to win a chess game against a human player. We chose to develop an AI specifically for Chess because Chess is a perfect information game that is particularly well-suited to intelligence analysis. To elaborate on that point, Chess is an extremely complex mathematical game and employs an endless amount of strategy and depth. In order for a player to win, one needs to strategically plan as the opponent makes each move.

BACKGROUND & MOTIVATION

For those who may not be as familiar with the board game, Chess is a two-player strategy game played on an 8x8 grid. Each player begins with 16 pieces: 8 pawns, 2 knights, 2 bishops, 2 rooks, 1 queen, and 1 king. These pieces are set up at opposite sides of the board and the goal is to place the opponent's king under checkmate, meaning that the king is unable to get itself to a safe place where it is not being attacked (How to Play Chess: Rules and Basics, n.d.). In certain cases, a game of chess will end in a draw, also known as a stalemate, where neither player is able to checkmate the opponent. That being said, the goal of our Chess AI is to win games; it will only draw if necessary.

Our main goal for this final project was to explore intelligent machines. Due to the wide array of logic that is used in Chess, we can actively determine the intelligence of our AI by

analyzing the moves it is making. As stated above, Chess is a mathematical game that heavily employs strategy. What this means is that a player needs to find a comfortable balance between attacking the opponent and defending its own pieces, especially the king. In order to get to checkmate, each move needs to be well-calculated. Since each chess piece has a given point value, we can easily differentiate good moves from those that are bad; this creates the basis of our intelligence analysis.

PERFECT INFORMATION

A perfect information game is a game where all of the players know everything about what has occurred throughout the game. Chess is a perfect information game because both players can see all of the pieces on the board, so they have all of the knowledge needed to figure out how each piece got to its current game location. Perfect information games are useful environments to study many of the different algorithms and decision making techniques that are used in artificially intelligent systems because they provide environments where guesswork is not involved. For example, an aspect of poker is being able to guess to some extent the cards in your opponent's' hands, and this makes it very difficult to create an effective poker AI. Chess, being a perfect information game, doesn't have this problem. The AI has access to all of the information it needs to be highly competitive. We would like to expand on what we have learned in class and use the perfect information game of chess as a basis to study various types of AI related algorithms. This paper will go into detail about the work that was done to accomplish the goal of creating an intelligent chess agent.

DESIGNING OUR CHESS AI

The structure of our Chess AI was extremely important in order to build an effective chess agent. In order of evaluation, the chess AI must know the moves that it can make, the moves the opponent can make depending on its own moves, and then finally use heuristics and algorithms to determine which is the best move. The difficulty of making a good chess AI is that it has to know how to strategize and make moves similar to that of an advanced human player. When we think of how chess masters (people who play at the highest level of chess) play, they use multiple angles of attack and constantly strategize 5-6 moves ahead in order to win. However, for an AI to be able to do that, we have to write the heuristics so that the AI has logic to play the game effectively and then algorithms so that it can strategize ahead of the current board state and, hopefully, win against its opponent.

The main difference between a chess AI and a human playing the game is that the AI needs a set of rules that evaluate to a numerical value that it can rely on to make intelligent moves, whereas a human has the ability to think and reason and can apply their knowledge to choosing the move they think is best. In addition, given the quickly changing nature of a chess game, an intelligent chess agent must have the ability to determine its best move every turn. It may also have the ability to look more than one ply deep, and make moves that a human player may not see.

COMMON CHESS AI ALGORITHMS

Computers or robots with artificial intelligence employ many different algorithms, including, but not limited to, learning algorithms and search algorithms (Hartikka, 2017). For our Chess AI, we chose a search algorithm and will go into detail in the next few paragraphs to show the thinking process behind our decision.

Learning algorithms are a common implementation of artificial intelligence in a computer. These algorithms take input from their environment and use it to train their model and produce an output. In 1959, Arthur Samuel defined machine learning as the “field of study that gives computers the ability to learn without being explicitly programmed.” Once the inputs are put into the learning model, they are cross referenced with the expected output, and the model is changed accordingly to try to output closer to the expected result. The model is then trained with various inputs and expected outputs until it has “learned” how to perform the calculations correctly.

Another type of algorithm used in artificial intelligence is a search algorithm. These algorithms tend to utilize a cause-and-effect concept where the search considers each possible action available at a given moment (Mayefsky, n.d.). These algorithms, combined with a set of rules, evaluate which decision is most intelligent. The difference between search algorithms and learning algorithms is that search algorithms do not change their model over time.

These types of models can be used to build an intelligent chess agent. In particular, search algorithms fit very well with chess since they evaluate and return an intelligent decision by considering each possible action at a given moment. Given a chess environment, this means the algorithm will consider all possible moves at a given turn and decide which of the moves is the best. Coupled with a learning algorithm, the agent can evaluate if the moves are truly intelligent and change its model accordingly over time.

CHESSNUT

For the base of this project we decided to use a Github project called Chessnut, developed by Chris Gearhart et al. We decided to work with Chessnut because it provides a simple python

chess model and stores the state of the game in Forsyth Edwards Notation (FEN). FEN was important to the implementation of this project because it made it simple to parse the board state into our graphical user interface (GUI). In addition, we chose Chessnut because it is not an engine and lacks its own GUI, so it did not achieve any of the goals we were trying to accomplish with our project. The library's clean code and useful methods set a solid stage for us to integrate our AI chess agent and GUI. Lastly, in order to focus the majority of our efforts on the AI, we settled on a simple GUI encoded in American Standard Code for Information Interchange (ASCII), which models an 8x8 chess board, and represents pieces with letters, designating color by the letter's case.

MINIMAX ALGORITHM & HEURISTICS

We believe that basic chess heuristics in combination with the minimax algorithm will successfully implement an intelligent chess agent. The role of the heuristics is to assign a numerical value to each of the chess pieces and subsequently each chess move has an associated point value. The minimax algorithm, by keeping track of the point values of the different pieces, maximizes the point value of the AI's potential move while minimizing the point value of the opponent's potential move based on the AI's potential move.

We chose this approach because it fits best with the manner in which chess is played. Specifically, having the minimax algorithm evaluate the best move at each turn should work well with the game of chess given the dynamic nature of the game. Although other approaches were considered, such as deep learning and reinforcement learning, we believed that it would be less effective to train a model on the best move since chess changes drastically each turn. This is why

we believed that it was best to implement a model that could take any board state and return the best move solely based on this information.

Minimax is a very well known algorithm that is used in many different contexts. It is best known for its application in two-player, zero-sum games (Two-Person Zero-Sum Games, n.d.). The algorithm is also used in game theory, a branch of mathematics used to evaluate competitive situations between two players. It is commonly used in business, finance, and economics, and relies on minimax to determine the best possible outcome. The aim of minimax in financial and business applications is to cause “the least amount of regret, should the strategy fail” (Minimax principle, n.d.).

Heuristics are also very popular and used in many contexts. A heuristic is viewed as a rule of thumb, or method of granting a short decision-making time based on a set of rules (Cherry, 2016). They not only apply to the field of computer science, but to psychology as well. In computer science, heuristics are used for both algorithms and user interface design. User interface design relies heavily on 10 heuristics that were developed by Jakob Nielsen, such as error prevention, user control and freedom among others (Nielsen, 1995). In general, heuristics are very commonly used as rules of thumb for quickly gaining intelligence on a situation.




Figure 1: Visualization of Minimax (Thornton, n.d.)

Looking at the figure above, minimax works by assigning alternating tree depths to moves between the player and the AI. For example, level 0 refers to the first move the AI would make and level 1 shows all the moves the player could make depending on the the AI's first move. The diagram is simplified for readability and in real life chess games level 1 would have many more moves and the tree would be much, much larger. After this, depending on the chosen ply or depth in the minimax algorithm, we will calculate points by looking at the root nodes and calculating a difference between the total points the AI accumulated versus those of the user. The highest positive number indicates the best move the AI can make, whereas the most negative number indicates the worst move for the AI and the best move for the opponent. After finding the best move for the AI, we propagate that value back up the tree and pass that as our best move.

Both minimax and heuristics were informed by psychological theories. Minimax has its base in social psychology, specifically the theory of social situations. The theory of social situations deals with the tradeoff of maximizing and minimizing between two individuals. The relationship between this theory and minimax is prevalent in the exchange between two agents, and their attempt to maximize for themselves, and minimize the other. Heuristics is also directly related to psychology, in that they are used in decision making and problem-solving. They are the way in which people quickly evaluate situations, though they may not necessarily be conscious of it. As pointed out by Kendra Cherry in her article "What Is a Heuristic and How Does It Work?", although heuristics are necessary for decision making, they can also lead to cognitive bias. She explains that "just because something has worked in the past does not mean that it will work again, and relying on an existing heuristic can make it difficult to see alternative

solutions or come up with new ideas.” This not only applies to heuristics in a psychological setting, but also to computer science. If heuristics are implemented in a program, it is important to realize that just because they worked before does not mean they will always be as effective. It was important to take this into account when designing our heuristics for the chess AI, because we had to make them general enough to apply to any game of chess, but specific enough to be effective.

PROJECT PLANNING

Our project was planned in predetermined iterations to simplify keeping track of progress and integrating new functionality. The first step was to create a GUI that worked with Chessnut to completely model a game of chess and be capable of taking in user input. The next iteration was to build in a simple AI. This AI chose randomly from a list of its next possible moves. Although this iteration was very simple, we believed it was an important landmark to make certain that we could build in a chess agent with Chessnut that could play against a human. After this was done, our next iteration was to implement heuristics into the agent so it could intelligently choose its next move.

The first heuristic we implemented was taking a random move. Though this was valuable for our development process because it meant we could play against an AI, we wanted to build an AI with better logic so that it had a better chance of winning games.

For our next iteration, we ranked moves by the point value of the pieces. At the start of the game, many, if not all, of the spaces in rows 3-6 of the chess board will be empty and those point values will be 0. However, once the game has gone into full swing, a pawn that has a possible move of taking a knight will be assigned a point value of 3 points. This demonstrated

that our AI had the capability of being intelligent, though basic, and pose a slight challenge to the human.

The last iteration was the most important for us to implement, though most difficult, which was to successfully integrate the minimax algorithm into our chess AI. This posed a huge challenge and took a lot of work to complete. After completion, though, our chess agent functions intelligently and provides fair competition to the human.

GROUPWORK

We divided our group work by looking at each of our team members' strengths and worked on different parts accordingly. Henry worked on coding the minimax algorithm and heuristics, building the node class, writing helper functions for the minimax algorithm such as determining what move was selected by minimax, and writing the readme file in markdown. Alex's strength was writing and he put together an in-depth outline for our final report and helped Henry with the minimax algorithm. Elias was tasked with implementing the GUI and integrating it with Chessnut, writing a function to run the game, building in the ability to look ahead 2 ply, and helping implement minimax.

Given that our group was a disadvantage and lacking a fourth person, we tried our best to distribute tasks and outline the progression of the project. Elias worked with Chessnut because he felt most comfortable reading through their API. Given his understanding of Chessnut, he helped the group work through all of the problems that we encountered with the library. Henry felt comfortable learning about the minimax algorithm and thus he wrote and implemented it into the code base. Alex voiced his strengths with writing and took charge of the outline, which led to the start of our paper.

Overall, we all helped each other out when we got stuck and worked as a cohesive group to move the project forward effectively. Decisions were always made as a group. The scope and outline of our iterations were decided as a group, and we monitored the rate at which work was completed, as to not fall behind.

CHALLENGES

One challenge we faced was conceptually understanding how the minimax algorithm worked in relation to a chess game. This difficulty in understanding the concept of minimax in regards to chess made it difficult to plan and implement the 2-ply solution. Chessnut also made the implementation of the minimax algorithm difficult due to the complexity of implementing a lookahead function in this environment. Unfortunately, Chessnut is built in a way that makes it very difficult to look ahead more than 1 ply. In order to look ahead two ply, a player must simulate a move and then generate a list of the 1 ply moves as a result of the initial move. Chessnut prevents a player from performing a move if it is not their turn, so figuring out how to lookahead 2 ply after a player's turn was difficult and involved learning how to parse and manipulate FEN. Given that a lookahead function is a necessary component of a chess AI, we were slowed down while we tried to figure out how to correctly implement the function.

RESULTS

We successfully created a chess agent that has the ability to intelligently choose its moves. Using different heuristics and algorithms, such as minimax, the chess agent can intelligently choose from its possible moves which is the best given the current state of the board.

Tests were done on the runtime and piece capturing ability of three algorithms, 1 Ply Random Move, 1 Ply Best Move, and 2 Ply Minimax. The results of the tests are below:

| Move Duration of Different Algorithms (seconds) | | | |
|--|-------------------|-----------------|---------------|
| | 1 Ply Random Move | 1 Ply Best Move | 2 Ply Minimax |
| Turn 1 Time | 9.060E-06 | 1.335E-02 | 5.752 |
| Turn 2 Time | 6.914E-06 | 1.535E-02 | 4.956 |
| Turn 3 Time | 7.868E-06 | 1.636E-02 | 16.383 |
| Turn 4 Time | 7.868E-06 | 1.656E-02 | 14.788 |
| Turn 5 Time | 8.821E-06 | 1.754E-02 | 16.234 |
| Turn 6 Time | 7.153E-06 | 1.622E-02 | 12.114 |
| Turn 7 Time | 8.106E-06 | 1.668E-02 | 15.100 |
| Turn 8 Time | 7.153E-06 | 1.512E-02 | 15.096 |
| Turn 9 Time | 8.106E-06 | 1.633E-02 | 16.190 |
| Turn 10 Time | 8.106E-06 | 1.528E-02 | 16.327 |
| Average Time | 7.915E-06 | 1.588E-02 | 13.294 |

Figure 2: Move Duration of Different Algorithms

The table above displays how long it takes the particular algorithms to make their moves. It can be seen that 1 Ply Random Move is the quickest, followed by 1 Ply Best Move, and 2 Ply Minimax. It should be noted that 2 Ply Minimax is significantly slower than the two other algorithms because it has to do many more computations to look two moves ahead for every possible chess move. 2 Ply Minimax is roughly 1.6 million times slower than 1 Ply Random Move, and 800 times slower than 1 Ply Best Move.

| Game Score of First 20 Moves (points) | Game 1 | Game 2 | Game 3 |
|--|--------|--------|--------|
| 2 Ply Minimax vs 1 Ply Random Move | | | |
| 2 Ply Minimax Score | 14 | 21 | 23 |
| 1 Ply Random Move Score | 5 | 0 | 9 |
| 2 Ply Minimax vs 1 Ply Best Move | | | |
| 2 Ply Minimax Score | 22 | 23 | 12 |
| 1 Ply Best Move Score | 9 | 15 | 9 |
| 1 Ply Random Move vs 1 Ply Best Move | | | |
| 1 Ply Random Move Score | 3 | 1 | 16 |
| 1 Ply Best Move Score | 10 | 22 | 18 |

Figure 3: Game Score of First 20 Moves

In a series of three games, the scores of different algorithms playing each other can be seen in the table over. In both games that 2 Ply Minimax played, it consistently outperformed the other algorithms. Also, it can be seen that 1 Ply Best Move performed better than 1 Ply Random Move in every case as well. From this test, 2 Ply Minimax is the best followed by 1 Ply Best Move with 1 Ply Random Move being the worst.

Overall, 2 Ply Minimax is the best algorithm because, although it is computationally slower, it delivers a significantly higher score than the other two algorithms. For each AI move, it looks two moves ahead to calculate a point differential and that is the key to continually having a point advantage over the two other agents which only look at the next best move. The logic behind minimax is made to be better than our other agents since the ply/depth directly corresponds with the number of moves it looks ahead. Thus, if it had a depth of 6, it would probably be able to beat most human players since people cannot keep track of the computations for all possible moves six turns down the road with their memories. The only downside to minimax is that as depth increases, performance decreases because the computation time increases exponentially with each additional depth.

One way to combat this is by adding Alpha-Beta pruning, which our team did not have time to implement. Pruning is harder to implement with chess compared to simpler two player games because our point value differential and list of possible moves is much larger and more complicated than typical games. An important thing to remember is that chess has six different pieces that all move in unique patterns. For example, a similar game, such as checkers, has one unique piece with one moving pattern.

NEXT STEPS

If we were to continue working on this project, we would first make it so that the minimax algorithm is functioning at a much greater depth, like 6-ply. We would then continue to improve our searching capabilities by adding more algorithms that speed up the search process such as the aforementioned Alpha-Beta pruning. Simultaneously, we would continue to add heuristics that make the AI perform better, such as the heuristic of trying to get pieces to control the center of the board or trying to castle as soon as possible.

CONCLUSION

This project served as a real-world example of how many of the topics, theories, and algorithms we have learned about throughout the semester are used in the field of AI Cognitive Science. Specifically in terms of chess, we learned about the established general approach to starting the construction of a chess-playing AI. As we conducted more research and continued to implement our solution, we saw how one could build off of basic chess AI by adding in the heuristics that are second nature to chess masters in order to create one of the celebrated chess AI's that can defeat world champions, for example. When we learned about the sheer number of search streamlining techniques and evaluation heuristics that go into making one of these master AIs, the team could not help but be humbled by the amount of effort and detail that goes into making the complex systems.

REFERENCES

- Cherry, K. (2016, June 20). What Is a Heuristic and How Does It Work? Retrieved April 25, 2017, from <https://www.verywell.com/what-is-a-heuristic-2795235>
- Chess Programming Wiki. (n.d.). Retrieved April 12, 2017, from <https://chessprogramming.wikispaces.com/>
- Hartikka, L. (2017, March 30). A step-by-step guide to building a simple chess AI – freeCodeCamp. Retrieved April 17, 2017, from <https://medium.freecodecamp.com/simple-chess-ai-step-by-step-1d55a9266977>
- How to Play Chess: Rules and Basics. (n.d.). Retrieved May 1, 2017, from <https://www.chess.com/learn-how-to-play-chess>
- Mayefsky, E., Anene, F., & Sirota, M. (n.d.). Retrieved May 02, 2017, from <http://www-cs-faculty.stanford.edu/people/eroberts/courses/soco/projects/2003-04/intelligent-search/minimax.html>
- Minimax principle. (n.d.). Retrieved April 24, 2017, from <http://www.businessdictionary.com/definition/minimax-principle.html>
- Nielsen, J. (1995, January 1). Nielsen Norman Group. Retrieved May 1, 2017, from <https://www.nngroup.com/articles/ten-usability-heuristics/>
- Thornton, C. (n.d.). KR-IST - Lecture 5a Game playing with Minimax and Pruning. Retrieved April 24, 2017, from <http://users.sussex.ac.uk/~christ/crs/kr-ist/lec05a.html>
- Two-Person Zero-Sum Games: Basic Concepts. (n.d.). Retrieved April 23, 2017, from <https://neos-guide.org/content/game-theory-basics>